

Single-cell foundation models represent genes in context, but reveal no context-specific gene function beyond co-expression

Ihor Kendiukhov · 2026

2026

Institute of Medical Genetics and Applied Genomics, University of Tübingen, Calwerstraße 7, 72076 Tübingen, Germany. Correspondence: kendiukhov@gmail.com

Abstract

Background: Single-cell foundation models read a cell as a sequence of its genes and build, at every layer, a representation of each gene that depends on the other genes in the cell. Whether this contextual representation carries information about a gene’s role in a given cell type, the single-cell counterpart of word sense in language models, has not been tested.

Results: We measured contextual gene representations in MaxToki (two sizes), scGPT and STATE, in 12 Tabula Sapiens cell types and an independent bone-marrow dataset. A gene’s representation changes with cell type in a gene-specific way: the same gene’s change agrees across independent cells at cosine 0.76, against 0.001 for a gene-shuffled control. The effect is strongest in the first two layers, survives rank control (99%), replicates on the independent dataset (0.71), is partly learned (untrained model: 0.44 against 0.74 on the same cells) and does not grow consistently with model size. Genes that change more gain somewhat more from real context in next-gene prediction (partial correlation 0.16). The changes follow functional axes such as nuclear versus surface genes ($z = 17$ against random axes), but gene sets matched on size and co-expression give similar organisation ($p = 0.08$ – 0.38). Steering along the nuclear/surface axis, with large pushes, shifts the model’s predictions for other genes ($p = 0.0081$); a second axis has no consistent effect and a third acts in reverse. Two tests for context-specific gene function, run in MaxToki-217M, were negative: transcription factors sit slightly closer to their curated targets than co-expression predicts, but not significantly more so where those targets are active. Functional organisation differs sharply between models ($z = 8.1$ – 18.4 for MaxToki, 3.2 for scGPT, 1.1 for STATE).

Conclusions: Single-cell foundation models build real contextual gene representations. In MaxToki, training organises them along functional axes no more than co-expression does, and we found no context-specific gene function beyond co-expression. Reading a gene’s changed role in a new cell type will likely need models trained on more than expression. The functional organisation of contextual change offers a simple metric for comparing architectures.

Keywords: Single-cell foundation models; Contextual gene representations; Mechanistic interpretability; Gene co-expression; Gene embeddings; Transformers; Single-cell RNA sequencing; Model evaluation

Background

A recurring theme in the interpretability of large language models is that a network trained on a narrow predictive task builds internal representations of structure it was never taught. Sequence models of Othello and chess build representations of the board and of latent variables such as

player skill (Li et al. 2023; Nanda et al. 2023; Karvonen 2024); language models linearly encode space and time along identifiable directions (Gurnee and Tegmark 2024); and, in the best cases, these representations are not decorative but causally used, a claim one tests by intervening on the representation and measuring the effect on behaviour (Turner et al. 2023; Zou et al. 2023; Park et al. 2024).

A closely related line concerns contextual representations. Contextual word embeddings assign the same word a different vector in every sentence, and the geometry of those vectors encodes word sense: “bank” near a river and “bank” in finance occupy separable regions (Peters et al. 2018; Ethayarajh 2019; Coenen et al. 2019). These measurements need care, because representation spaces are strongly anisotropic (vectors cluster in a narrow cone rather than spreading evenly, so any two look similar) and a handful of “rogue” dimensions can dominate cosine similarity (Ethayarajh 2019; Mu et al. 2018; Cai et al. 2021; Timkey and van Schijndel 2021).

Single-cell foundation models (SCFMs) offer a clean biological version of the same question (Theodoris et al. 2023; Cui et al. 2024; Hao et al. 2024; Yang et al. 2022). They treat each gene as a token and each cell as a set or sequence of genes, either ordered by expression (Geneformer) or paired with binned or continuous expression values (scGPT, scBERT, scFoundation), and are trained to predict masked or next genes or values, in effect learning which genes tend to be active together. Their gene token embeddings (one fixed vector per gene) have been studied (Cui et al. 2024; Kendiukhov 2026g); their contextual gene representations (a gene’s vector inside a particular cell) are used as features downstream (Hao et al. 2024; Bian et al. 2024) but have not been characterised as objects in their own right. Yet this is exactly where a model could, in principle, represent that a gene plays a different role in a macrophage than in a T cell: the biological analogue of word sense, and an open direction that the “cell as a token” framing invites but does not test (Gilpin 2025).

Related work. Our methods descend directly from the language-model literature on contextualised embeddings. Ethayarajh (Ethayarajh 2019) showed that contextual representations are anisotropic in all non-input layers, that same-word self-similarity falls in upper layers, and that every cosine must be corrected against an anisotropy baseline. Timkey and van Schijndel (Timkey and van Schijndel 2021) showed that a few rogue dimensions dominate similarity and that a scalar correction is not enough; Mu et al. (Mu et al. 2018) and Cai et al. (Cai et al. 2021) give the standard postprocessing and the local-isotropy picture. Coenen et al. (Coenen et al. 2019) and Garí Soler and Apidianaki (Garí Soler and Apidianaki 2021) show that word senses form separable clusters and that a word’s degree of polysemy is itself readable. We import both the measurements and the controls. From the literature on emergent, causally used latent variables (Li et al. 2023; Nanda et al. 2023; Karvonen 2024; Gurnee and Tegmark 2024) we adopt the evidentiary standard that a probe shows only correlation, so a causal claim needs an intervention against a matched control, together with its steering methods (Turner et al. 2023; Zou et al. 2023), of which the extraction and amplification of interpretable feature directions in production models is the large-scale instance (Templeton et al. 2024).

Geneformer (Theodoris et al. 2023), scGPT (Cui et al. 2024), scFoundation (Hao et al. 2024), scBERT (Yang et al. 2022), UCE (Rosen et al. 2026) and STATE (Adduri et al. 2026) differ in tokenisation (rank-ordered gene identity, binned or continuous expression values, or gene tokens derived from protein-sequence embeddings), and the review of Bian et al. (Bian et al. 2024) defines the “contextual gene embedding” that our work dissects. MaxToki (Gómez Ortega et al. 2026), the primary model here, is a Llama-architecture, Geneformer-lineage model at 217M and 1B scales. Two earlier analyses come closest to a “context changes the gene” claim. scGPT reports cell-state-specific gene networks read from attention (Cui et al. 2024), but as a pairwise network, not as a property of the representation vector, and without a decomposition or a null. Geneformer’s in silico deletion measures how the embeddings of the other genes shift when one gene is removed from a cell (Theodoris et al. 2023), which is a change of context by one gene. Neither characterises how a gene’s representation differs between natural cell types. The geometry of gene representations in

these models has been mapped, both for static gene embeddings (Kendiukhov 2026g) and for layer-wise hidden states averaged over cells, including their spectral structure (Kendiukhov 2026c) and the topological and geometric structure they do and do not encode, per tissue (Kendiukhov 2026h). Residual-stream gene vectors built separately in several tissues carry a small, tissue-dependent amount of transcription-factor-target signal beyond expression covariates (Kendiukhov 2026d). How much, and how gene-specifically, a single gene’s vector changes from one cell type to another, the object here, has not been measured.

A critical literature finds that SCFMs often fail to beat simple baselines: on zero-shot tasks (Kedzierska et al. 2025), under deeper evaluation (Boiarsky et al. 2024), on perturbation prediction (Ahlmann-Eltze et al. 2025; Bendidi et al. 2024), and across utilities (Liu et al. 2026). Sparse autoencoders find organised biological knowledge but minimal regulatory logic in these models (Kendiukhov 2026e), and edge scores read from SCFM attention add nothing beyond simple gene-level statistics for predicting which genes respond to a perturbation (Kendiukhov 2026f). Co-expression is the canonical guilt-by-association signal (van Dam et al. 2018), and some gene-embedding methods, such as Gene2vec, are built directly on it (Du et al. 2019); others build gene embeddings from text descriptions of genes instead (Chen and Zou 2025). Genuine context-dependent gene function exists. Some of it acts at the protein level, as in moonlighting proteins (Jeffery 1999), and leaves no expression signature. Some acts through transcription, as when lineage-determining transcription factors select different enhancers with different partners by cell type (Heinz et al. 2010).

Aim. We organise the question into two levels. Level 1 (a fact about the model): is a gene’s representation reshaped by context in a structured, gene-specific way, and is that reshaping learned and used? Level 2 (biology): does the reshaping reveal something about the gene’s role in each context beyond what co-expression already tells us? We answer Level 1 for MaxToki-217M with a rank control, an independent replication, a random-weights control, a link to the model’s predictions and a steering test, and compare contextualisation strength and functional organisation across four models. For Level 2, we test the functional organisation of the contextualisation against co-expression nulls and run two direct tests for context-specific gene function, all in MaxToki-217M. The study makes three contributions: (1) the Level 1 characterisation; (2) the Level 2 boundary; and (3) contextualisation as a model-characterisation metric.

Methods

Models

MaxToki-217M and MaxToki-1B (Gómez Ortega et al. 2026; Theodoris Lab 2026) are Llama-architecture autoregressive models over rank-ordered gene tokens, with 11 and 20 transformer layers and hidden widths of 1232 and 2304. We used the Hugging Face safetensors checkpoints through the `transformers` library with `output_hidden_states=True`, in float32. For MaxToki, “layer k ” is the k -th hidden state: layer 0 is the embedding output and layer $k - 1$ is the output of the k -th transformer block; the last hidden state (layer 11 of MaxToki-217M) is taken after the model’s final normalisation. scGPT (Cui et al. 2024; Bo Wang Lab, n.d.) (whole-human checkpoint; 12 layers; width 512) and STATE (Adduri et al. 2026; Arc Institute 2025) (the SE-600M state-embedding model; 16 layers; width 2048) were used for the cross-architecture comparison. STATE’s gene tokens are fixed protein embeddings from the Evolutionary Scale Modeling 2 (ESM2) language model (Lin et al. 2023), passed through a learned projection together with each gene’s expression. For scGPT and STATE, “layer k ” is the output of the transformer module with index k (the $(k+1)$ -th block). As an untrained control, we built a `LlamaForCausalLM` from MaxToki-217M’s configuration with random initial weights (the library’s default initialisation, seed 0) and the same vocabulary.

Datasets and annotations

The main data are the Tabula Sapiens (Tabula Sapiens Consortium 2022) immune, kidney and lung datasets from CZ CELLxGENE Discover (CZI Cell Science Program et al. 2025; Tabula Sapiens Consortium, n.d.-a, n.d.-b, n.d.-c). For the immune dataset we used a 20,000-cell subset; its cell identifiers are listed in the code repository. These files store log-normalised expression ($\log_{1p}$ of counts per 10,000) in their main matrix and the original integer counts in their raw matrix. All MaxToki analyses and the co-expression matrix read the raw counts; the scGPT and STATE inputs are described below. Contexts are cell types: the 12 largest cell types with at least 250 cells (macrophage, pulmonary alveolar type 2 cell, kidney epithelial cell, capillary endothelial cell, CD4-positive alpha-beta T cell, CD8-positive alpha-beta T cell, B cell, basal cell, pulmonary alveolar type 1 cell, intermediate monocyte, neutrophil and classical monocyte), with 1000 randomly chosen cells per cell type.

For independent replication we used all 5780 cells of the human CD34+ bone-marrow data of Setty et al. (Setty et al. 2019) as distributed with scVelo (Bergen et al. 2020; scVelo project, n.d.), which holds velocityto spliced and unspliced integer counts for 14,319 pre-filtered genes (11,842 in the MaxToki vocabulary); its main matrix holds the spliced counts only. Because the Tabula Sapiens counts include intronic reads, the main Setty analysis ranks genes by the sum of spliced and unspliced counts; a robustness run used spliced counts only. Contexts are the nine differentiation clusters with at least 130 cells: two haematopoietic stem cell (HSC) clusters, precursors, common lymphoid progenitors (CLP), two erythroid clusters, two monocyte clusters and dendritic cells (DC). Megakaryocytes (61 cells) were excluded. All cells of each kept cluster were used.

Gene Ontology (GO) annotations (Ashburner et al. 2000; Gene Ontology Consortium et al. 2023) were taken from the gene-to-GO file distributed with GEARS (Roohani et al. 2024, 2022). A gene belongs to a functional pole only if it is annotated directly to one of the listed GO terms; annotations were not propagated through the GO hierarchy. Curated transcription-factor (TF) targets were all human TRRUST v2 interactions (Han et al. 2018b, 2018a) together with DoRothEA interactions of confidence A or B (Garcia-Alonso et al. 2019; Saez-Rodriguez Lab, n.d.). Gene symbols were mapped to Ensembl identifiers with the Geneformer dictionaries (version gc104M) (Theodoris et al. 2023; Theodoris, n.d.).

Input encoding and extraction of contextual gene representations

MaxToki uses Geneformer’s rank-value encoding (Theodoris et al. 2023): a cell’s counts are normalised to the cell total, each gene is divided by its median value from the model’s gene-median dictionary, and the non-zero genes are ranked in decreasing order. A per-cell scale factor does not change the order within a cell, so the order is that of raw count divided by gene median. The first 1022 genes were kept and wrapped in the beginning- and end-of-sequence tokens (1024 tokens at most). On real cells, our implementation gives the same token sequences as the MaxToki tokenizer, except that genes with exactly tied count-to-median values can appear in swapped order.

We recorded, for each gene in a panel, its hidden state at that gene’s token position and averaged it over the cells of one cell type. The main MaxToki-217M extraction used layers 0, 4, 8 and 11. A separate extraction of all 12 hidden states (the embedding output and all 11 transformer layers) with a 4000-gene panel, on the same cells, gave the full depth profile. Four design choices each fix a failure we measured in a pilot. (1) Pairwise contexts, not the all-context intersection: in the main extraction only 137 genes are count-balanced in all 12 cell types, many of them broadly expressed housekeeping genes (the single-cell analogue of the “stopwords” that are most context-sensitive in language models (Ethayarajh 2019)), whereas each pair of cell types shares a median of 1391 count-balanced genes. The panel contains genes with at least 25 occurrences in at least two cell types (the 6000 most frequent such genes). (2) Cell-level split halves: each cell was assigned to one of two partitions, so every per-gene mean is estimated twice from disjoint cells and gives an honest noise floor. (3) An occurrence cap: a per-gene mean has a standard error proportional to $1/\sqrt{n}$, and

unequal counts inject context-dependent noise into any interaction term, so at most 50 occurrences were accumulated per (gene, cell type, partition), over a shuffled cell order. A (gene, cell type) entry is used only if it reached the cap in both partitions (“count-balanced”). (4) Per-dimension z-scoring: before any cosine, every hidden dimension was standardised over all count-balanced entries, to remove the inflation by a few dominant (“rogue”) dimensions that a scalar anisotropy correction misses (Timkey and van Schijndel 2021).

Other extractions used the same procedure with these settings. Scaling pair: MaxToki-217M and MaxToki-1B on the same 600 cells per cell type with a 5000-gene panel (cap 50), at layers 2 and 4 of MaxToki-217M and layers 4 and 7 of MaxToki-1B, so that the models can be compared at matched relative depth (layer 4 of 11 = 0.36 versus layer 7 of 20 = 0.35; layer 2 of 11 = 0.18 versus layer 4 of 20 = 0.20). Random-weights control: the same cells and the same 5000-gene panel as the MaxToki-217M scaling run, at layers 4 and 8. Cap-matched control: the first 20 occurrences of each (gene, cell type, partition) in the main extraction’s forward pass, at layer 4. Setty data: cap 25, a floor of 15 occurrences in at least two clusters, a 5000-gene panel, layers 2 and 4. scGPT: representations from an existing per-gene activation atlas of 3000 Tabula Sapiens cells (immune, kidney and lung) (Kendiukhov 2026e), in which each cell’s 1200 highest-expressed genes were given to scGPT with their log-normalised values passed to the value channel without scGPT’s standard binning; we reformatted it into the same layout: cell types with at least 100 cells (8 cell types), two partitions, cap 20, floor 12, 5000 genes. STATE: a preprocessed copy of each tissue that keeps about 4800 highly variable genes, with ambient-corrected counts; 600 cells per tissue (immune, kidney and lung), taken as up to 150 cells from each of the largest cell types until 600 were reached (in practice the 4 largest per tissue); STATE fills its input with unexpressed genes as well, and we kept only positions holding expressed genes (count > 0); cell types with at least 100 cells (9 cell types), cap 20, floor 12 (1840 genes).

Two-way decomposition and contextualisation strength (EXCESS)

For gene g in cell type c , we write $v(g, c) = \mu + a(g) + b(c) + E(g, c)$: a gene main effect $a(g)$ (the gene’s identity), a context main effect $b(c)$ (a shift shared by all genes in that cell type, expected because attention pools over the cell), and a gene-specific interaction $E(g, c)$, the only term where a gene-specific response to context can appear.

For a cell-type pair (c_0, c_1) , the relative shift of gene g is $\delta(g) = [v(g, c_1) - v(g, c_0)] - \text{mean}_g [v(g, c_1) - v(g, c_0)]$: the gene’s own shift minus the average shift of all genes. We computed δ separately in the two cell partitions (c_0, c_1) and defined $\text{EXCESS} = \text{mean cos}(\delta_0(g), \delta_1(g)) - \text{mean cos}(\delta_0(g), \delta_1(\text{perm } g))$: the same gene’s shift agreeing across independent cells, minus a gene-shuffled control. $\text{EXCESS} = 0$ means no gene-specific response to context. The means are pooled over all (gene, pair) entries from the cell-type pairs with at least 200 count-balanced genes (150 in the cross-model and Setty analyses), which equals a mean over pairs weighted by their number of genes. 95% confidence intervals (CIs) come from 2000 bootstrap resamples of genes, in which each resampled gene brings all of its pair entries. As a positive control we report the replication of the context main effect across partitions (the cosine between the two partitions’ mean shifts).

Anisotropy. On the raw (unstandardised) representations we computed the mean cosine of random different-gene pairs (the anisotropy baseline (Ethayarajh 2019)), same-gene self-similarity across cell types, and the share of variance on the first principal component (PC), and repeated these after per-dimension z-scoring.

Rank-position control. MaxToki encodes token position, and a gene’s rank changes between cell types. We computed each gene’s mean rank in each cell type over the cells that express it, using the same cells as the extraction. We split genes at the median absolute rank change into rank-stable and rank-moving halves and recomputed EXCESS in each, correlated the size of the shift with the rank change, and recomputed EXCESS after regressing the shift on the gene’s mean rank in each of the two cell types and on its absolute rank change.

Functional axes and reproducible interaction power

A functional axis u is the unit vector separating two GO-defined gene classes in the gene main effect, $u = \text{mean}(a(g) \mid \text{class A}) - \text{mean}(a(g) \mid \text{class B})$, where $a(g)$ is the gene’s z-scored representation averaged over partitions and count-balanced cell types. Genes annotated to both poles were removed from both. The axes were nuclear versus surface/secreted (nucleus GO:0005634, chromatin GO:0000785 and DNA binding GO:0003677, versus plasma membrane GO:0005886, extracellular region GO:0005576 and extracellular space GO:0005615); mitochondrion (GO:0005739) versus cytoskeleton (GO:0005856); transcription (regulation of DNA-templated transcription GO:0006355 and DNA-binding transcription factor activity GO:0003700) versus transport and signalling (ion transport GO:0006811 and signalling receptor activity GO:0038023); and cell cycle (GO:0007049) versus cell differentiation (GO:0030154). Axis validity is the five-fold cross-validated area under the receiver operating characteristic (ROC) curve (AUC) with which $a(g)$, u separates held-out pole genes.

Reproducible interaction power measures how much context moves genes along u in a way that replicates across independent cells. For genes count-balanced in at least 9 cell types, we projected each partition’s representation onto u , $q_p(g, c) = v_p(g, c) \cdot u$; regressed out the gene’s mean rank in that cell type (linear and quadratic terms; “rank-controlled”); took the two-way interaction $I_p(g, c)$ (the projection minus its gene mean and context mean, plus the grand mean); and averaged $I_0(g, c) \cdot I_1(g, c)$ over all (gene, cell type) entries. Noise averages to zero in this product, so only modulation that reproduces across the two partitions survives. Functional-z is the z-score of this power against 300 axes built from random gene partitions of the same pole sizes. We report it at layers 4, 8 and 11, without and with the rank regression. A cell type’s loading on an axis is the mean projection onto u of the z-scored layer-4 representations (averaged over partitions) of the genes count-balanced in that cell type. These gene sets differ between cell types and the loadings are not rank-controlled, so we use them only to describe the ordering of cell types.

Null models for functional organisation

Functional gene sets may be co-expressed, and co-expressed genes co-vary across cells whatever the model “understands”. We therefore compared each functional axis’s power with null axes whose two poles have the same sizes as the functional poles. Co-expression was the Pearson correlation of $\log_{10}(\text{counts per } 10,000)$, computed from the raw counts over 3000 cells from each of the three Tabula Sapiens datasets, and the coherence of a gene set was its mean pairwise correlation.

1. **Random axes:** random genes (300 axes).
2. **Strongly co-expressed modules:** each null pole was the set of genes most correlated with a random seed gene (300 axes). These modules are much more co-expressed than the functional poles (Results), so this is a strong null rather than a matched one.
3. **Size- and co-expression-matched sets:** each null pole was built from the k genes most correlated with a random seed gene plus randomly chosen genes, with k chosen by bisection so that the pole’s coherence was within 10% of the functional pole’s (300 axes). We also traced power against the fraction of seed-co-expressed genes at fixed pole size (0 to 100%, 25 draws per step).
4. **Size- and tightness-matched sets:** the same construction with representation tightness (the mean pairwise cosine of the pole genes’ main effects $a(g)$) in place of co-expression (300 axes).

For each null, the empirical one-sided p is $(k + 1)/(N + 1)$, where k is the number of the $N = 300$ null axes with power at least as large as the functional axis; when no null axis reached it, p takes its smallest possible value, $1/301$. Empirical p values were used because the null power distributions are right-skewed.

Causal steering test

To test whether the model uses the functional direction, we added it to the residual stream at half of a cell’s gene positions and read the model’s own next-gene output logits at the other half. The direction was the raw (unstandardised) difference between the mean layer-4 representations of the nuclear-pole and surface-pole genes (poles made disjoint, as for the axes). It was added at the layer-4 representation (the output of the fourth transformer block). For each of 30 immune cells (at most 512 tokens), gene positions were randomly split into a steered half and a read half, so any effect at the read positions must travel through attention. The readout was the mean change in logit of nuclear-pole gene tokens minus that of surface-pole gene tokens (specificity). Strengths were 0.25, 0.5 and 1.0 times the mean residual-stream norm at that layer in one reference cell; we report 0.5 as the headline. Pushing toward $+u$ and toward $-u$ and taking the difference (the signed swing) cancels any effect that does not depend on the sign of the push. The control was a random direction of the same norm, drawn fresh for each cell. Significance was a one-sided sign test of functional swing greater than random swing across cells. The test was repeated at the layer-8 representation and for the mitochondrion/cytoskeleton and transcription/transport axes at layer 4. To put the push sizes in context, we measured how far context naturally moves genes along the same direction: the spread of each gene’s projection onto u across its count-balanced cell types.

Tests for context-specific gene function

Directional congruence. Along each functional axis, we decomposed each gene’s projection into a gene loading $f(g)$ (mean over contexts), a context loading $h(c)$ (mean over genes) and an interaction $e(g, c)$. If genes move toward their own functional pole specifically in cell types that share that character, e should follow $f \cdot h$. We estimated f and h in one partition and e in the other (and vice versa), measured $\text{corr}(e, f \cdot h)$ on rank-residualised projections, and compared with 300 null axes whose poles were matched to the functional poles on size and co-expression (as in null 3 above).

Placement near curated targets. For TFs count-balanced in at least 6 cell types with at least 15 curated targets in the gene panel, we asked (A) whether, in each cell type, the TF’s representation (with the context mean removed) is closer by cosine to the centroid of its curated targets than to the centroid of non-targets matched on co-expression with the TF (40 redraws), and (B) whether this excess closeness rises in cell types where the targets are more highly expressed (Spearman correlation across cell types with target activity, the negative of the targets’ mean rank). A cell type was used for a TF only if at least 8 of its targets were count-balanced there, and a TF needed at least 3 such cell types. Both questions were tested across TFs with sign tests.

Lineage transcription factors. We also attempted a test on 14 TFs with lineage-specific roles in immune cells (SPI1, IRF4, IRF8, GATA3, MAFB, BCL6, PRDM1, TBX21, RUNX1, RUNX3, BATF, MAF, BACH2 and ZEB2) (Heinz et al. 2010), but few of them were sampled well enough in the panel (Results).

Cross-model, scaling and random-weights comparison

EXCESS and nuclear/surface functional-z were computed identically for every model at layer 4, and for the scaling pair also at matched relative depth, on each model’s own gene vocabulary (context pairs with at least 150 count-balanced genes). In this comparison, and in the Setty replication, functional-z was computed without the rank regression, on genes count-balanced in at least $n/2$ of the n cell types (4 of 9 for STATE and the Setty clusters), against 200 random-partition axes. Because the occurrence cap differs between architectures (50 for MaxToki, 20 for scGPT and STATE), we also report MaxToki-217M with cap 20.

Link to next-gene prediction

For 350 Tabula Sapiens immune cells with at least 60 expressed vocabulary genes (at most 512 tokens), and up to 20 target positions per cell from the sixth gene onward, we compared the model’s log-probability of the true gene given the cell’s own preceding genes with that given the same number of preceding genes from another cell (a chimeric context). The donor cell was drawn at random among cells with at least as many genes whose preceding genes did not include the target gene, so that the chimeric context never repeats the target. The per-gene context benefit (log-probability difference, averaged over occurrences; genes with at least 3 occurrences) was correlated with the gene’s per-gene context movement, the mean over its count-balanced cell types (at least 4) of $\|z(g, c) - \text{mean}_c z(g, c)\|$ in the z-scored layer-4 space (a per-gene quantity, distinct from EXCESS). We used a partial Spearman correlation controlling for log gene frequency, with a 95% CI from 2000 bootstrap resamples of genes.

Statistics and reproducibility

All analyses are deterministic (random seed 0) and run on cached extractions. Empirical p values are one-sided unless stated. Multiple tests were corrected within each family: the four steering configurations and the four curated-target tests (static closeness and its modulation, at layers 4 and 8) with a Bonferroni correction, and the six directional-congruence tests (three axes, layers 4 and 8; only positive counted) with a Holm correction. Each analysis has its own script: input encoding `ctx_tokenise.py`; extraction `ctx_extract_maxtoki.py` (and `_scgpt`, `_state`, `_devel`, `_random`); `ctx_polysemy.py`, `ctx_layer_curve.py` and `ctx_anisotropy.py` (contextualisation by layer and anisotropy); `ctx_position_confound.py` (rank control); `ctx_independent.py` (Setty replication); `ctx_functional_axes.py`, `ctx_celltype_loadings.py`, `ctx_coexpr_null_v2.py`, `ctx_coexpr_null_v3.py` and `ctx_tightness_null.py` (functional axes and nulls); `ctx_causal.py` (steering); `ctx_directional_probe.py`, `ctx_curated_targets.py` and `ctx_switcher_test.py` (tests for context-specific function); `ctx_cross_model.py` (cross-model comparison); `ctx_prediction_link.py` (prediction link); and `ctx_figures.py` (figures). Each writes a result file `results/ctx_*.json`, from which every number, table and figure in this paper is taken. The runs used Python 3.12 with PyTorch 2.12.1, transformers 5.13.0 and NumPy 2.4 on Apple-silicon hardware.

Use of large language models

An artificial intelligence (AI) assistant (Claude, Anthropic) was used to help write and run the analysis code, check references, and draft and edit the manuscript text. The author designed the study, reviewed the code, results and text, and takes full responsibility for the content. The assistant is not an author.

Results

All analyses below use MaxToki-217M unless a model is named. We first establish Level 1 (contextualisation is real, not rank position, replicable, partly learned and weakly linked to prediction, and one functional direction is read by the model when pushed), then report the Level 2 tests.

Genes are contextualised gene-specifically, most strongly in the first layers

The gene-specific context response is large and highly reproducible (Table 1, Fig. 1). At layer 4 of the main extraction, EXCESS is +0.758: the same gene’s relative shift agrees at cosine +0.759 across independent cells, while a different gene’s agrees at +0.001, so the effect is gene-specific. The context main effect replicates at +0.996, which shows that the measurement works. Layer 0 (the context-free embedding output) returns exactly 0.000, the check that the pipeline reports nothing when there is nothing. The full depth scan shows where the response arises: it is highest at layer

1 (+0.847) and layer 2 (+0.835) and then falls to +0.651 at layer 11 (Fig. 1). Context is therefore integrated in the first one or two transformer layers. Deeper layers keep most of it (90% of the peak at layer 4 and 84% at layer 8), while the gene-specific shift becomes somewhat less reproducible toward the output.

Table 1: **Table 1.** Gene-specific context response at the four layers of the main MaxToki-217M extraction

Layer	EXCESS (same – different)	Same gene	Different gene	Main-effect replication
0 (embedding)	+0.000	+0.000	+0.000	—
4	+0.758	+0.759	+0.001	+0.996
8	+0.711	+0.711	+0.000	+0.996
11	+0.649	+0.650	+0.001	+0.997

EXCESS is the agreement of the same gene’s relative context shift across two independent halves of the cells, minus the agreement for a different gene, pooled over 66 cell-type pairs (a median of 1391 count-balanced genes per pair). Gene-level bootstrap 95% CIs are within ± 0.003 . The full depth profile is shown in Fig. 1.

Metric validity (anisotropy). These representations are anisotropic, as expected (Ethayarajh 2019): the mean cosine of random different-gene pairs is +0.20 to +0.62 at layers 4–11 (raw). Same-gene self-similarity across cell types falls with depth (+0.92 $\rightarrow$ +0.86 $\rightarrow$ +0.69 at layers 4, 8 and 11), so genes move more between cell types in upper layers, the single-cell echo of Ethayarajh’s finding. Together with the EXCESS profile, this means that upper layers move genes more but less reproducibly. Per-dimension z-scoring brings anisotropy to about 0.00 while keeping genes distinguishable (corrected self-similarity +0.88/+0.65/+0.62), and the first principal component carries only 0.07 of the variance at layer 11, so no single dimension dominates the cosines.

The effect is not the gene’s rank position

MaxToki reads a cell as a rank-ordered gene list and encodes token position, so a gene ranking 5th in one cell type and 500th in another is an obvious confound. It is not the driver (Fig. 2). The gene-specific response is about the same for rank-stable genes (+0.749) and rank-moving genes (+0.767) at layer 4, and 99% of it survives regression on rank (+0.749); 98% and 97% survive at layers 8 and 11. The size of a gene’s shift correlates only weakly with its rank change ($r = 0.26$).

The effect replicates on an independent dataset

On the Setty CD34+ bone-marrow data, the lineage in which a recoverable manifold was found in scGPT (Kendiukhov 2026a), with nine differentiation clusters rather than adult cell types as contexts, the phenomenon replicates: EXCESS is +0.713 at layer 2 and +0.713 at layer 4 (gene-shuffled null +0.0036 and +0.0039), and functional-z is +5.0 and +5.0 for the nuclear/surface axis and +16.1 and +8.8 for the mitochondrion/cytoskeleton axis. Using spliced counts only gives the same picture (EXCESS +0.702 and nuclear/surface functional-z +9.3 at layer 4). Neither the contextualisation nor its functional organisation is specific to Tabula Sapiens or to adult cell types.

Context moves genes along functional axes, as co-expression does

Context moves genes along interpretable functional axes far more than along random directions (Table 2). At layer 4, functional-z (rank-controlled in parentheses) is +16.6 (+13.3) for nuclear versus surface/secreted genes, +8.7 (+6.4) for mitochondrion versus cytoskeleton and +2.9 (+4.3) for transcription versus transport. The cell-cycle versus differentiation axis, which separates its poles

only weakly (validity AUC 0.57), gives +2.0 (+0.9). The organisation persists in deeper layers for the two strongest axes. Cell types order sensibly along the nuclear/surface axis: lymphocytes sit at the nuclear end (CD8 T +4.7, CD4 T +4.4, B +3.9) and alveolar epithelial cells and macrophages at the surface end (alveolar type 1 −3.6, type 2 −3.2, macrophage −2.8).

Table 2: **Table 2.** Functional organisation of context modulation, by axis and layer (MaxToki-217M)

Axis	Poles (genes)	Validity AUC (layer 4)	Layer 4	Layer 8	Layer 11
Nuclear / surface	1800/1188	0.77	+16.6 (+13.3)	+7.8 (+6.5)	+6.1 (+6.0)
Mitochondrion / cytoskeleton	675/165	0.78	+8.7 (+6.4)	+12.7 (+8.4)	+10.4 (+6.9)
Transcription / transport	326/89	0.69	+2.9 (+4.3)	+1.0 (+1.3)	+0.8 (+0.8)
Cell cycle / differentiation	145/167	0.57	+2.0 (+0.9)	+2.9 (+1.7)	+2.6 (+1.9)

Functional-z is the z-score of the axis’s reproducible interaction power against 300 axes built from random gene sets of the same pole sizes, without (and, in parentheses, with) regression on the gene’s rank. Validity AUC is the cross-validated area under the ROC curve with which the axis separates held-out pole genes.

The key control asks whether this organisation exceeds co-expression (Fig. 3). Functional gene sets can be co-regulated, and co-regulated genes co-vary across cells whatever the model “understands”. The functional poles are more co-expressed than random genes, but far less than co-expression modules (mean pairwise correlation 0.028 and 0.042 for the nuclear and surface poles, against 0.025 for random gene sets). Against null gene sets matched on size and this average co-expression, no functional axis exceeds the null: the nuclear/surface power (0.50) sits inside the matched null (mean 0.40, 95th percentile 0.68; $p = 0.316$), and $p = 0.375$ and 0.083 for the other two axes. Modules built from the genes most co-expressed with a seed gene, 3–11 times more co-expressed than the functional poles, reach higher power on average (nuclear/surface null mean 0.67; $p = 0.701$, 0.757 and 0.804). At fixed pole size, 66 of 175 null axes with 0–100% seed-co-expressed genes reached the nuclear/surface power. At layer 4 the organisation is not a by-product of compact gene sets in representation space: it exceeds sets matched on size and representation tightness ($p = 0.007$, 0.023 and 0.013). At layer 8 it does not ($p = 0.233$, 0.060 and 0.100). It far exceeds random axes (none of 300 reached it; $p = 1/301$ for nuclear/surface). So the functional organisation of the contextualisation is no stronger than that of gene sets with the same co-expression.

The nuclear/surface direction is read by the model, but not every functional direction is

In the steering test (Fig. 4), adding the nuclear/surface direction to half of a cell’s genes moves the model’s own logits at the untouched half with the sign of the push: nuclear-gene logits rise for a push toward the nuclear pole (specificity +0.085) and fall for a push toward the surface pole (−0.097), a signed swing of +0.182 against +0.001 for a random direction of the same norm. The swing is +0.079, +0.182, +0.184 at strengths 0.25, 0.5 and 1.0 and exceeds the random control in 22 of 30 cells ($p = 0.0081$; Bonferroni-adjusted 0.032). The effect is stronger at the layer-8 representation (swing +0.232, 30 of 30 cells, $p = 9.3 \times 10^{-10}$). The direction therefore propagates through attention and changes the model’s predictions.

This does not generalise to every functional axis. The mitochondrion/cytoskeleton direction has no

significant effect at strengths 0.25 and 0.5 (swing +0.024, +0.067, +0.269; at 0.5, 15 of 30 cells, $p = 0.57$); at strength 1.0 it exceeds the random control in 20 of 30 cells ($p = 0.049$, Bonferroni-adjusted 0.20), the same as the nuclear/surface direction at that strength. The transcription/transport direction has a reversed effect that grows with strength (swing -0.139 , -0.268 , -0.678): the functional swing exceeds the random one in only 2 of 30 cells (two-sided sign test $p = 8.7 \times 10^{-7}$). We have no explanation for the reversal; it shows that an axis that separates genes is not automatically read in the expected direction. The pushes are also large compared with natural variation: the middle strength (7.6 units) is about 3 times the median range of a gene’s projection on the nuclear/surface direction across cell types (2.62). The test therefore shows that the model can read this axis, not that the natural, context-driven changes along it are used.

No context-specific gene function beyond co-expression

Two direct tests for context-appropriate gene function find no context-specific signal beyond co-expression. (1) Directional congruence: genes do not move toward their own functional pole specifically in cell types that share that character more than matched null axes do. At layer 4, μ is +0.16, +0.10 and +0.23 for the three axes, all positive but within the upper tail of size- and co-expression-matched null axes, whose means are near zero ($p = 0.116$, 0.086 and 0.053; none significant after Holm correction over 6 tests; at layer 8, μ is -0.16 , -0.11 and +0.09). The null spread of μ (standard deviation 0.08–0.13) means the test could detect only clear congruence. (2) Placement near curated targets: across 26 TFs, at layer 4 a TF’s representation is slightly closer to the centroid of its curated targets than to that of co-expression-matched non-targets (mean cosine excess +0.021; 20 of 26 TFs positive; sign test $p = 0.005$, Bonferroni-adjusted over 4 tests 0.019). This closeness does not rise significantly in the cell types where the targets are more active (mean $\mu = +0.073$; 17 of 26 TFs positive; $p = 0.08$, Bonferroni-adjusted 0.34), and at layer 8 neither effect is significant (static $p = 0.279$; modulation $\mu = +0.135$, $p = 0.08$). So the model places TFs modestly near their curated targets beyond what co-expression with the TF predicts, but not in a context-appropriate way. A third attempt, on 14 lineage TFs, is uninformative rather than negative: only 1 of them (RUNX1) was sampled well enough, because these regulators are weakly expressed in the cells we used.

Contextualisation as a model metric: architecture, scale, and learned versus architectural

Because it does not depend on resolving gene function, contextualisation can serve as a model metric (Table 3, Fig. 5). Three findings stand out. (1) Every model contextualises genes strongly (EXCESS 0.71–0.83), but contextualisation does not grow consistently with model size. At matched relative depth, MaxToki-1B gives +0.767 (layer 7) against +0.740 for MaxToki-217M (layer 4), and +0.767 (layer 4) against +0.831 (layer 2). Functional organisation is not higher in the larger model either (functional-z +8.1 against +14.3, and +18.4 against +18.4). (2) Functional organisation differs sharply between architectures: the rank-based MaxToki models give $z = 8.1$ –18.4 (Table 3), scGPT +3.2 and STATE +1.1. STATE moves genes strongly in context (EXCESS +0.828) but in no functionally coherent direction; its gene tokens are fixed protein embeddings rather than learned from co-expression. The models differ in more than tokenisation (training data, objective, gene panel and context sets), so this comparison does not isolate the cause. (3) The phenomenon is partly architectural and partly learned. An untrained MaxToki of the same architecture, run on the same cells and gene panel, keeps about 60% of the trained contextualisation (EXCESS +0.441 against +0.740) but only about a third of its functional organisation (functional-z +4.7 against +14.3). Fixed random embeddings and attention mixing give a reproducible, gene-specific shift on their own; training strengthens it and aligns it with functional axes. The untrained MaxToki’s functional-z (+4.7) is higher than that of trained scGPT (+3.2) and STATE (+1.1), so part of what this metric measures comes from MaxToki’s rank-based input and architecture rather than from training.

Table 3: **Table 3.** Contextualisation strength and functional organisation across models

Model	Layer (relative depth)	Contexts	Genes	Width	EXCESS	Functional- z
scGPT	4 (0.42)	8	5000	512	+0.705	+3.2
STATE	4 (0.31)	9	1840	2048	+0.828	+1.1
(SE-600M)						
MaxToki- 217M	2 (0.18)	12	5000	1232	+0.831	+18.4
MaxToki- 217M	4 (0.36)	12	5000	1232	+0.740	+14.3
MaxToki- 1B	4 (0.20)	12	5000	2304	+0.767	+18.4
MaxToki- 1B	7 (0.35)	12	5000	2304	+0.767	+8.1
MaxToki- 217M, random weights	4 (0.36)	12	5000	1232	+0.441	+4.7

EXCESS measures contextualisation strength; functional-z measures organisation along the nuclear/surface axis against 200 random-gene axes (no rank regression). Relative depth is the number of transformer blocks applied divided by the model’s total number of blocks: for MaxToki, layer k follows block k ; for scGPT and STATE, layer 4 follows the fifth block (5/12 and 5/16). The MaxToki rows use the same 600 cells per cell type and the same 5000-gene panel (cap 50); scGPT and STATE come from their own extractions (cap 20), so their EXCESS should not be ranked against MaxToki’s (Limitations).

Contextualisation is weakly aligned with what the model uses for prediction

If contextualisation serves the autoregressive objective, the genes the model contextualises most should be those whose next-gene prediction gains most from real context. That is what we find. Real context helps prediction a great deal: across 2266 comparisons for 558 genes, the mean benefit is +12.0 nats (median +12.9; a nat is a natural-log unit of probability), so the true gene is roughly 160,000-fold more probable under its real context than under another cell’s. The per-gene size of this benefit tracks how much the gene moves between cell types: the partial Spearman correlation, controlling for gene frequency, is +0.161 (95% CI +0.083 to +0.239). Genes whose representations depend more on context tend, weakly, to be genes whose prediction depends more on context.

Discussion

The results characterise how SCFMs represent genes in context and what that representation carries. On the positive side, MaxToki reshapes gene representations gene-specifically and reproducibly, most strongly in its first one or two layers. Training strengthens the reshaping (an untrained model keeps about 60% of its strength and a third of its functional organisation); the reshaping is weakly aligned, gene by gene, with how much the model’s predictions depend on context; and in the MaxToki models it follows functional axes. To our knowledge this is the first measurement for single-cell models of the “how contextual are contextualised representations” question (Ethayarajh 2019), and the comparison of models gives an interpretability metric that separates the three architectures we tested, though not the two MaxToki sizes consistently.

What the reshaping carries is, as far as our tests can tell, co-expression. The functional organisation of the context changes is no stronger than that of gene sets matched to the functional sets on size and average co-expression, and modules of strongly co-expressed genes produce more; we did not test directly whether a representation built from expression alone reproduces it. The two direct tests for context-specific gene function are negative. The one signal beyond co-expression is static: TFs sit slightly closer to their curated targets than to co-expression-matched non-targets at layer 4, but not more so in the cell types where those targets are active. The steering result fits the same picture. The model reads the nuclear/surface direction, but the pushes tested (3.8–15.2 units) are 1.4 to 6 times a gene’s typical natural range along it, so the test does not show that natural context changes are used; and two other functional directions are not consistently read, or are read in reverse. The architecture comparison does not contradict this: STATE, whose gene tokens are fixed protein embeddings, contextualises genes strongly but without functional organisation, and scGPT, also trained on expression, shows little; the comparison cannot separate tokenisation from the other differences between the models.

One possible reason for this boundary is that an expression-only model sees only which genes are expressed and their rank order, and its objective rewards predicting co-expressed genes. The alignment between contextualisation and the benefit of real context for prediction is consistent with that reading, although the test measures only how far a gene moves, not in which direction, so it does not show that the model moves genes toward their co-expressed partners. Some context-dependent function acts at the protein level, as in moonlighting proteins (Jeffery 1999), and leaves no expression signature. Other forms, such as cell-type-specific enhancer choice by lineage-determining TFs (Heinz et al. 2010), act through target-gene expression and could in principle be visible; our tests did not detect them. Whether multi-omic, sequence-aware or larger expression models reach Level 2 is not tested here.

Limitations

Most analyses were run on MaxToki-217M only: the rank control, the Setty replication, the steering test, the co-expression and tightness nulls, the prediction link and both Level 2 tests. MaxToki-1B, scGPT and STATE were tested only for contextualisation strength and nuclear/surface functional organisation. Cross-architecture EXCESS is confounded by the occurrence cap (20 for scGPT and STATE versus 50 for MaxToki) and by width. A MaxToki-217M run with the same cap of 20 gives EXCESS +0.620 and functional-z +28.0, so the functional-z gap between architectures is not caused by the cap, but EXCESS values should not be ranked across architectures. The functional-z comparison uses non-identical context sets (12, 9 and 8 cell types) and, for STATE, a smaller gene panel (1840 genes) and only 14 of 36 cell-type pairs with enough count-balanced genes; we did not re-run it on a common context set. STATE was given only about 4800 highly variable genes per tissue (ambient-corrected counts), which is likely one reason its panel holds only 1840 genes and few cell-type pairs could be scored. The rank control was run only for MaxToki; scGPT and STATE also receive each gene’s expression level, so part of their EXCESS may reflect the gene’s own expression change rather than its context. The scGPT representations come from an existing extraction in which log-normalised values were given to scGPT’s value channel without its standard binning, so its absolute activations are off its training distribution; our analyses use contrasts that largely cancel this, but the scGPT numbers should be read as relative to MaxToki, not as calibrated values. The scaling comparison has two model sizes only.

Most analyses use layer 4, not the peak layers 1–2; the co-expression nulls were run only at layer 4 and the tightness null at layers 4 and 8. The steering direction is the context-averaged nuclear-minus-surface direction, a gene-class direction rather than the context-specific change itself, and it was compared only with a random push, not with co-expression-matched directions. The co-expression nulls match average pairwise co-expression; power also depends on how co-expression is distributed within a gene set, and we did not run the more direct test of whether a representation built only from expression (for example, a factorisation of the same cells) reproduces the functional organisation.

The Setty data cover a pre-filtered set of genes, unlike the full raw counts of Tabula Sapiens. The Level 2 negatives are evidence of a ceiling, not proof of absence: some context-dependent function is invisible to an expression-only model, and the directional test could detect only clear congruence. Functional axes are built from direct annotations in one annotation file; the conclusions hold across three axes, but specific z-scores depend on how the poles are defined.

Conclusions

Single-cell foundation models build contextual gene representations that are a genuine object of study: a gene’s vector is reshaped by its cellular context in a gene-specific and reproducible way that training strengthens, most strongly in the first layers, weakly aligned with how much the model’s predictions rely on context, and different between architectures. To our knowledge, this is the first systematic characterisation of the phenomenon for single-cell models. The functional organisation of the contextual change gives a simple metric that separates the three architectures we tested. In MaxToki, controls matched on co-expression also bound what the contextualisation carries: its functional organisation is no stronger than that of gene sets with the same co-expression, and two direct tests for context-specific gene function are negative. Reading a gene’s changed role in a new cell type, the use that would make these representations a hypothesis-generation tool, is not reachable with the expression-only model tested here and will likely require models trained on more than expression.

List of abbreviations

AI: artificial intelligence; AUC: area under the receiver operating characteristic curve; CI: confidence interval; CLP: common lymphoid progenitor; DC: dendritic cell; ESM2: Evolutionary Scale Modeling 2; GO: Gene Ontology; HSC: haematopoietic stem cell; PC: principal component; ROC: receiver operating characteristic; SCFM: single-cell foundation model; TF: transcription factor.

Declarations

Ethics approval and consent to participate

Ethics approval and consent to participate were not required for this study, which is a secondary analysis of publicly available, de-identified human single-cell RNA-sequencing data (Tabula Sapiens (Tabula Sapiens Consortium 2022) and Setty et al. (Setty et al. 2019)) and involved no contact with participants. The original studies obtained ethics approval and informed consent, as described in those publications. No new human data or samples were collected.

Consent for publication

Not applicable.

Availability of data and materials

The datasets analysed during the current study are available in the following public repositories: the Tabula Sapiens immune, kidney and lung datasets in CZ CELLxGENE Discover (Tabula Sapiens Consortium, n.d.-a, n.d.-b, n.d.-c); the human CD34+ bone-marrow data of Setty et al. (Setty et al. 2019), as distributed with scVeloc, in figshare (scVeloc project, n.d.); the MaxToki checkpoints in Hugging Face (Theodoris Lab 2026); the scGPT whole-human checkpoint (Bo Wang Lab, n.d.); the STATE SE-600M model in Hugging Face (Arc Institute 2025); the Geneformer gene dictionaries in Hugging Face (Theodoris, n.d.); the GEARS gene-to-GO annotation file in Harvard Dataverse (Roohani et al. 2022); TRRUST v2 (Han et al. 2018a); and DoRothEA (Saez-Rodriguez Lab, n.d.). All analysis code, the result files from which every number, table and figure is computed, the

figure code and the list of the 20,000 immune cells used are available in the gene-context repository (Kendiukhov 2026b). The MaxToki, random-weights, Setty and STATE representation tensors (about 9 GB) are not deposited because of their size; they can be regenerated from the public inputs with the provided extraction scripts and are available from the corresponding author on reasonable request. The scGPT representations were reformatted from an activation atlas produced in earlier work (Kendiukhov 2026e); that atlas is available from the corresponding author on reasonable request. No new experimental data were generated.

Software availability:

- Project name: gene-context
- Project home page: <https://github.com/Biodyn-AI/gene-context>
- Archived version: git tag `bmc-submission-v1` of the project repository
- Operating system(s): platform independent (developed and run on macOS, Apple silicon)
- Programming language: Python 3.12
- Other requirements: NumPy, SciPy, scikit-learn, h5py and matplotlib; for the extraction scripts, PyTorch (2.2 or later) and transformers (4.43 or later), and for STATE the Arc Institute `state` package and `anndata`; versions used: PyTorch 2.12.1, transformers 5.13.0, NumPy 2.4
- License: MIT
- Any restrictions to use by non-academics: none

Competing interests

The author declares that he has no competing interests.

Funding

Not applicable. This research received no specific grant from any funding agency in the public, commercial or not-for-profit sectors.

Authors' contributions

IK conceived and designed the study, wrote the software, carried out the analyses, interpreted the results and wrote the manuscript. The author read and approved the final manuscript.

Acknowledgements

The author thanks the Tabula Sapiens Consortium, the authors of the Setty et al. bone-marrow study, and the developers of MaxToki, scGPT, STATE, Geneformer, GEARS, TRRUST and DoRothEA for making their data, models and resources publicly available.

References

- Adduri, Abhinav K., Dhruv Gautam, Beatrice Bevilacqua, et al. 2026. "Predicting Cellular Responses to Perturbation Across Diverse Contexts with State." *Cell* 189 (19): 5914–5931.e20. <https://doi.org/10.1016/j.cell.2026.07.052>.
- Ahlmann-Eltze, Constantin, Wolfgang Huber, and Simon Anders. 2025. "Deep-Learning-Based Gene Perturbation Effect Prediction Does Not yet Outperform Simple Linear Baselines." *Nature Methods* 22 (8): 1657–61. <https://doi.org/10.1038/s41592-025-02772-6>.
- Arc Institute. 2025. "State Embedding Model SE-600M." Hugging Face. <https://huggingface.co/arcinstitute/SE-600M>.

- Ashburner, Michael, Catherine A. Ball, Judith A. Blake, et al. 2000. “Gene Ontology: Tool for the Unification of Biology.” *Nature Genetics* 25 (1): 25–29. <https://doi.org/10.1038/75556>.
- Bendidi, Ihab, Shawn Whitfield, Kian Kenyon-Dean, et al. 2024. “Benchmarking Transcriptomics Foundation Models for Perturbation Analysis: One PCA Still Rules Them All.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2410.13956>.
- Bergen, Volker, Marius Lange, Stefan Peidli, F. Alexander Wolf, and Fabian J. Theis. 2020. “Generalizing RNA Velocity to Transient Cell States Through Dynamical Modeling.” *Nature Biotechnology* 38 (12): 1408–14. <https://doi.org/10.1038/s41587-020-0591-3>.
- Bian, Haiyang, Yixin Chen, Erpai Luo, et al. 2024. “General-Purpose Pre-Trained Large Cellular Models for Single-Cell Transcriptomics.” *National Science Review* 11 (11): nwae340. <https://doi.org/10.1093/nsr/nwae340>.
- Bo Wang Lab. n.d. “scGPT Source Code and Pretrained Whole-Human Checkpoint.” GitHub. Accessed October 1, 2026. <https://github.com/bowang-lab/scGPT>.
- Boiarsky, Rebecca, Nalini M. Singh, Alejandro Buendia, Ava P. Amini, Gad Getz, and David Sontag. 2024. “Deeper Evaluation of a Single-Cell Foundation Model.” *Nature Machine Intelligence* 6 (12): 1443–46. <https://doi.org/10.1038/s42256-024-00949-w>.
- Cai, Xingyu, Jiaji Huang, Yuchen Bian, and Kenneth Church. 2021. “Isotropy in the Contextual Embedding Space: Clusters and Manifolds.” International Conference on Learning Representations (ICLR). <https://openreview.net/forum?id=xYGNO86OWDH>.
- Chen, Yiqun, and James Zou. 2025. “Simple and Effective Embedding Model for Single-Cell Biology Built from ChatGPT.” *Nature Biomedical Engineering* 9: 483–93. <https://doi.org/10.1038/s41551-024-01284-6>.
- Coenen, Andy, Emily Reif, Ann Yuan, et al. 2019. “Visualizing and Measuring the Geometry of BERT.” *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*.
- Cui, Haotian, Chloe Wang, Hassaan Maan, et al. 2024. “scGPT: Toward Building a Foundation Model for Single-Cell Multi-Omics Using Generative AI.” *Nature Methods* 21 (8): 1470–80. <https://doi.org/10.1038/s41592-024-02201-0>.
- CZI Cell Science Program, Shibli Abdulla, Brian Aevermann, et al. 2025. “CZ CELLxGENE Discover: A Single-Cell Data Platform for Scalable Exploration, Analysis and Modeling of Aggregated Data.” *Nucleic Acids Research* 53 (D1): D886–900. <https://doi.org/10.1093/nar/gkae1142>.
- Du, Jingcheng, Peilin Jia, Yulin Dai, Cui Tao, Zhongming Zhao, and Degui Zhi. 2019. “Gene2vec: Distributed Representation of Genes Based on Co-Expression.” *BMC Genomics* 20 (Suppl 1): 82. <https://doi.org/10.1186/s12864-018-5370-x>.
- Ethayarajh, Kawin. 2019. “How Contextual Are Contextualized Word Representations? Comparing the Geometry of BERT, ELMo, and GPT-2 Embeddings.” *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)* (Hong Kong), 55–65. <https://doi.org/10.18653/v1/D19-1006>.
- Garcia-Alonso, Luz, Christian H. Holland, Mahmoud M. Ibrahim, Denes Turei, and Julio Saez-

- Rodriguez. 2019. “Benchmark and Integration of Resources for the Estimation of Human Transcription Factor Activities.” *Genome Research* 29 (8): 1363–75. <https://doi.org/10.1101/gr.240663.118>.
- Garí Soler, Aina, and Marianna Apidianaki. 2021. “Let’s Play Mono-Poly: BERT Can Reveal Words’ Polysemy Level and Partitionability into Senses.” *Transactions of the Association for Computational Linguistics* 9: 825–44. https://doi.org/10.1162/tacl_a_00400.
- Gene Ontology Consortium, Suzi A. Aleksander, James Balhoff, et al. 2023. “The Gene Ontology Knowledgebase in 2023.” *Genetics* 224 (1): iyad031. <https://doi.org/10.1093/genetics/iyad031>.
- Gilpin, William. 2025. “The Cell as a Token: High-Dimensional Geometry in Language Models and Cell Embeddings.” *Bioinformatics* 41 (11): btaf595. <https://doi.org/10.1093/bioinformatics/btaf595>.
- Gómez Ortega, Javier, Rangarajan D. Nadadur, Akira Kunitomi, et al. 2026. “Temporal AI Model Predicts Drivers of Cell State Trajectories Across Human Aging.” *bioRxiv*, ahead of print. <https://doi.org/10.64898/2026.03.30.715396>.
- Gurnee, Wes, and Max Tegmark. 2024. “Language Models Represent Space and Time.” *The Twelfth International Conference on Learning Representations (ICLR 2024)*.
- Han, Heonjong, Jae-Won Cho, Sangyoung Lee, et al. 2018a. “TRRUST V2 Human Transcriptional Regulatory Interactions (Trrust_rawdata.human.tsv).” GRNpedia. <https://www.grnpedia.org/trrust/>.
- Han, Heonjong, Jae-Won Cho, Sangyoung Lee, et al. 2018b. “TRRUST V2: An Expanded Reference Database of Human and Mouse Transcriptional Regulatory Interactions.” *Nucleic Acids Research* 46 (D1): D380–86. <https://doi.org/10.1093/nar/gkx1013>.
- Hao, Minsheng, Jing Gong, Xin Zeng, et al. 2024. “Large-Scale Foundation Model on Single-Cell Transcriptomics.” *Nature Methods* 21 (8): 1481–91. <https://doi.org/10.1038/s41592-024-02305-7>.
- Heinz, Sven, Christopher Benner, Nathanael Spann, et al. 2010. “Simple Combinations of Lineage-Determining Transcription Factors Prime Cis-Regulatory Elements Required for Macrophage and B Cell Identities.” *Molecular Cell* 38 (4): 576–89. <https://doi.org/10.1016/j.molcel.2010.05.004>.
- Jeffery, Constance J. 1999. “Moonlighting Proteins.” *Trends in Biochemical Sciences* 24 (1): 8–11. [https://doi.org/10.1016/S0968-0004\(98\)01335-8](https://doi.org/10.1016/S0968-0004(98)01335-8).
- Karvonen, Adam. 2024. “Emergent World Models and Latent Variable Estimation in Chess-Playing Language Models.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2403.15498>.
- Kedzierska, Kasia Z., Lorin Crawford, Ava P. Amini, and Alex X. Lu. 2025. “Zero-Shot Evaluation Reveals Limitations of Single-Cell Foundation Models.” *Genome Biology* 26: 101. <https://doi.org/10.1186/s13059-025-03574-x>.
- Kendiukhov, Ihor. 2026a. “Discovery of a Hematopoietic Manifold in scGPT Yields a Method for Extracting Performant Algorithms from Biological Foundation Model Internals.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2603.10261>.

- Kendiukhov, Ihor. 2026b. “gene-context: Analysis Code, Result Files and Figures for This Study.” GitHub. <https://github.com/Biodyn-AI/gene-context>.
- Kendiukhov, Ihor. 2026c. “Multi-Dimensional Spectral Geometry of Biological Knowledge in Single-Cell Transformer Representations.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2602.22247>.
- Kendiukhov, Ihor. 2026d. “Residual-Stream Geometry of Single-Cell Foundation Models Carries Incremental Gene-Regulatory Signal Across Tissues.” *BMC Bioinformatics* 27 (1): 162. <https://doi.org/10.1186/s12859-026-06538-5>.
- Kendiukhov, Ihor. 2026e. “Sparse Autoencoders Reveal Organized Biological Knowledge but Minimal Regulatory Logic in Single-Cell Foundation Models: A Comparative Atlas of Geneformer and scGPT.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2603.02952>.
- Kendiukhov, Ihor. 2026f. “Systematic Evaluation of Single-Cell Foundation Model Interpretability: Attention-Derived Edge Scores Add No Incremental Value over Gene-Level Features for Perturbation-Target Prediction.” *BMC Genomics* 27 (1): 634. <https://doi.org/10.1186/s12864-026-12965-8>.
- Kendiukhov, Ihor. 2026g. “Two Axes in the Gene-Embedding Space of Single-Cell Foundation Models: What a Protein Language Model Already Knows, and What Requires Single-Cell Pre-training.” *SSRN*, ahead of print. <https://doi.org/10.2139/ssrn.6577900>.
- Kendiukhov, Ihor. 2026h. “What Topological and Geometric Structure Do Biological Foundation Models Learn? Evidence from 141 Hypotheses.” *PLOS One* 21 (7): e0344826. <https://doi.org/10.1371/journal.pone.0344826>.
- Li, Kenneth, Aspen K. Hopkins, David Bau, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. 2023. “Emergent World Representations: Exploring a Sequence Model Trained on a Synthetic Task.” *The Eleventh International Conference on Learning Representations (ICLR 2023)*.
- Lin, Zeming, Halil Akin, Roshan Rao, et al. 2023. “Evolutionary-Scale Prediction of Atomic-Level Protein Structure with a Language Model.” *Science* 379 (6637): 1123–30. <https://doi.org/10.1126/science.ade2574>.
- Liu, Tianyu, Kexing Li, Yuge Wang, Hongyu Li, and Hongyu Zhao. 2026. “Evaluating the Utilities of Foundation Models in Single-Cell Data Analysis.” *Advanced Science* 13 (27): e14490. <https://doi.org/10.1002/advs.202514490>.
- Mu, Jiaqi, Suma Bhat, and Pramod Viswanath. 2018. “All-but-the-Top: Simple and Effective Postprocessing for Word Representations.” *The Sixth International Conference on Learning Representations (ICLR 2018)*.
- Nanda, Neel, Andrew Lee, and Martin Wattenberg. 2023. “Emergent Linear Representations in World Models of Self-Supervised Sequence Models.” *Proceedings of the 6th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP* (Singapore), 16–30. <https://doi.org/10.18653/v1/2023.blackboxnlp-1.2>.
- Park, Kiho, Yo Joong Choe, and Victor Veitch. 2024. “The Linear Representation Hypothesis and the Geometry of Large Language Models.” *Proceedings of the 41st International Conference on Machine Learning. Proceedings of Machine Learning Research*, vol. 235, 39643–66.

- Peters, Matthew E., Mark Neumann, Mohit Iyyer, et al. 2018. “Deep Contextualized Word Representations.” *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)* (New Orleans), 2227–37. <https://doi.org/10.18653/v1/N18-1202>.
- Roohani, Yusuf, Kexin Huang, and Jure Leskovec. 2022. “GEARS Gene-to-Gene Ontology Annotation File (Gene2go.pkl; Harvard Dataverse File 6153417).” Harvard Dataverse. <https://dataverse.harvard.edu/api/access/datafile/6153417>.
- Roohani, Yusuf, Kexin Huang, and Jure Leskovec. 2024. “Predicting Transcriptional Outcomes of Novel Multigene Perturbations with GEARS.” *Nature Biotechnology* 42 (6): 927–35. <https://doi.org/10.1038/s41587-023-01905-6>.
- Rosen, Yanay, Yusuf Roohani, Ayush Agrawal, et al. 2026. “Universal Cell Embedding Provides a Foundation Model for Cell Biology.” *Nature* 656 (8126): 183–91. <https://doi.org/10.1038/s41586-026-10689-z>.
- Saez-Rodriguez Lab. n.d. “DoRothEA Human Regulons.” GitHub. Accessed October 1, 2026. <https://github.com/saezlab/dorothea>.
- scVelo project. n.d. “Human CD34+ Bone Marrow Data of Setty Et Al. As Distributed with scVelo (Bonemarrow.h5ad; Figshare File 27686835).” Figshare. Accessed October 1, 2026. <https://ndownloader.figshare.com/files/27686835>.
- Setty, Manu, Vaidotas Kisieliovas, Jacob Levine, Adam Gayoso, Linas Mazutis, and Dana Pe’er. 2019. “Characterization of Cell Fate Probabilities in Single-Cell Data with Palantir.” *Nature Biotechnology* 37 (4): 451–60. <https://doi.org/10.1038/s41587-019-0068-4>.
- Tabula Sapiens Consortium. 2022. “The Tabula Sapiens: A Multiple-Organ, Single-Cell Transcriptomic Atlas of Humans.” *Science* 376 (6594): eabl4896. <https://doi.org/10.1126/science.abl4896>.
- Tabula Sapiens Consortium. n.d.-a. “Tabula Sapiens – Immune (Dataset Version E62b0182-368f-4875-8ac1-F49e3d5eda60, Schema 7.0.0).” CZ CELLxGENE Discover. Accessed October 1, 2026. <https://datasets.cellxgene.cziscience.com/e62b0182-368f-4875-8ac1-f49e3d5eda60.h5ad>.
- Tabula Sapiens Consortium. n.d.-b. “Tabula Sapiens – Kidney (Dataset Version 65ca6e36-73b0-4c88-B0f3-7b23b48844ad, Schema 7.0.0).” CZ CELLxGENE Discover. Accessed October 1, 2026. <https://datasets.cellxgene.cziscience.com/65ca6e36-73b0-4c88-b0f3-7b23b48844ad.h5ad>.
- Tabula Sapiens Consortium. n.d.-c. “Tabula Sapiens – Lung (Dataset Version 40f8b1a3-9f76-4ac4-8761-32078555ed4e, Schema 7.0.0).” CZ CELLxGENE Discover. Accessed October 1, 2026. <https://datasets.cellxgene.cziscience.com/40f8b1a3-9f76-4ac4-8761-32078555ed4e.h5ad>.
- Templeton, Adly, Tom Conerly, Jonathan Marcus, et al. 2024. “Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet.” Transformer Circuits Thread. <https://transformer-circuits.pub/2024/scaling-monosemanticity/>.
- Theodoris, Christina V. n.d. “Geneformer: Gene Name, Gene Median and Token Dictionaries.” Hugging Face. Accessed October 1, 2026. <https://huggingface.co/ctheodoris/Geneformer>.
- Theodoris, Christina V., Ling Xiao, Anant Chopra, et al. 2023. “Transfer Learning Enables Predictions in Network Biology.” *Nature* 618 (7965): 616–24. <https://doi.org/10.1038/s41586-023-06139-9>.

- Theodoris Lab. 2026. “MaxToki Model Checkpoints (217M and 1B, Hugging Face Safetensors).” Hugging Face. <https://huggingface.co/theodoris-lab/MaxToki>.
- Timkey, William, and Marten van Schijndel. 2021. “All Bark and No Bite: Rogue Dimensions in Transformer Language Models Obscure Representational Quality.” *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing* (Punta Cana), 4527–46. <https://doi.org/10.18653/v1/2021.emnlp-main.372>.
- Turner, Alexander Matt, Lisa Thiergart, Gavin Leech, et al. 2023. “Steering Language Models with Activation Engineering.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2308.10248>.
- van Dam, Sipko, Urmo Vösa, Adriaan van der Graaf, Lude Franke, and João Pedro de Magalhães. 2018. “Gene Co-Expression Analysis for Functional Classification and Gene-Disease Predictions.” *Briefings in Bioinformatics* 19 (4): 575–92. <https://doi.org/10.1093/bib/bbw139>.
- Yang, Fan, Wenchuan Wang, Fang Wang, et al. 2022. “scBERT as a Large-Scale Pretrained Deep Language Model for Cell Type Annotation of Single-Cell RNA-seq Data.” *Nature Machine Intelligence* 4 (10): 852–66. <https://doi.org/10.1038/s42256-022-00534-z>.
- Zou, Andy, Long Phan, Sarah Chen, et al. 2023. “Representation Engineering: A Top-down Approach to AI Transparency.” *arXiv*, ahead of print. <https://doi.org/10.48550/arXiv.2310.01405>.

Figures




Figure 1: **Figure 1. Gene-specific context response from the embedding through all 11 layers of MaxToki-217M.** Agreement of the same gene’s relative context shift across independent halves of the cells (blue; full depth scan with a 4000-gene panel) against the gene-shuffled null (grey). Open diamonds mark the four layers of the main 6000-gene extraction (Table 1). Layer 0, the context-free embedding output, is exactly zero; the response is highest at layer 1 (+0.847) and falls to +0.651 at layer 11. Values are pooled over 66 cell-type pairs per layer; gene-level bootstrap 95% CIs are narrower than the markers.




Figure 2: **Figure 2. The gene-specific context response is not explained by rank position.** EXCESS at layers 4, 8 and 11 for genes whose rank changes little between cell types (rank-stable, below the median rank change), genes whose rank changes a lot (rank-moving, above the median), and all genes after regressing the shift on the gene's ranks and rank change. The response is about the same for rank-stable and rank-moving genes, and 99%, 98% and 97% of it survives the rank regression at layers 4, 8 and 11.




Figure 3: **Figure 3. Functional organisation compared with random and co-expression nulls.** For each of the three functional axes with well-separated poles (layer 4; the cell-cycle/differentiation axis is omitted), the red marker and line show the functional axis's rank-controlled reproducible interaction power. Boxes show 300 null axes each, all with the functional poles' sizes: random genes (grey); gene sets matched on size and average co-expression (light blue); and modules of the genes most co-expressed with a random seed gene (dark blue), which are 3–11 times more co-expressed than the functional poles. Boxes show the interquartile range and median; whiskers reach the most extreme null value within 1.5 times the interquartile range (values beyond are not drawn). Annotated p values are empirical one-sided values against the matched null and against the strongly co-expressed modules. The functional axes are far above random axes but within both co-expression nulls.




Figure 4: **Figure 4. The nuclear/surface direction is read by the model.** Signed swing (the change in nuclear-minus-surface logit specificity for a push toward the nuclear pole minus that for a push toward the surface pole) at the untouched half of each cell’s genes, when the nuclear/surface direction is added at the layer-4 representation of the other half, as a function of steering strength (multiples of the mean residual-stream norm). Blue: functional direction (mean $\pm$ standard error over 30 cells). Grey: random direction of the same norm. The mitochondrion/cytoskeleton direction has no significant effect at strengths 0.25 and 0.5, and the transcription/transport direction a reversed one (Results).




Figure 5: **Figure 5. Contextualisation as a model metric across architecture, scale and training.** Left: contextualisation strength (EXCESS). Right: functional organisation (nuclear/surface functional-z against random-gene axes, without rank regression); the dashed line marks $z = 3$, a reference level for a clearly non-random effect. Rows: scGPT and STATE (SE-600M) at layer 4, each from its own extraction (occurrence cap 20); MaxToki-217M at layer 4 and MaxToki-1B at layer 7 (matched relative depth, 0.36 and 0.35), on the same 600 cells per cell type (cap 50); and MaxToki-217M with random weights (grey). Every trained model contextualises strongly, but only the MaxToki models show strong functional organisation; random weights give weaker contextualisation and weak functional organisation. Because the caps differ, EXCESS should not be ranked across architectures (Limitations).